

Class 10: Halloween Candy Mini Project

Barry

Background

In this mini-project we will examine 538 Halloween Candy data. What is your favorite candy? What is nougat anyway? And how do you say it in America?

First step is to read the data...

```
candy <- read.csv("candy-data.csv", row.names=1)
head(candy)
```

	chocolate	fruity	caramel	peanutyalmondy	nougat	crispedricewafer
100 Grand	1	0	1	0	0	1
3 Musketeers	1	0	0	0	1	0
One dime	0	0	0	0	0	0
One quarter	0	0	0	0	0	0
Air Heads	0	1	0	0	0	0
Almond Joy	1	0	0	1	0	0

	hard	bar	pluribus	sugarpercent	pricepercent	winpercent
100 Grand	0	1	0	0.732	0.860	66.97173
3 Musketeers	0	1	0	0.604	0.511	67.60294
One dime	0	0	0	0.011	0.116	32.26109
One quarter	0	0	0	0.011	0.511	46.11650
Air Heads	0	0	0	0.906	0.511	52.34146
Almond Joy	0	1	0	0.465	0.767	50.34755

Q1. How many different candy types are in this dataset?

```
nrow(candy)
```

```
[1] 85
```

Q2. How many fruity candy types are in the dataset?

```
sum(candy$fruity)
```

```
[1] 38
```

```
candy["Twix", ]$winpercent
```

```
[1] 81.64291
```

```
rownames(candy)
```

```
[1] "100 Grand"           "3 Musketeers"
[3] "One dime"            "One quarter"
[5] "Air Heads"           "Almond Joy"
[7] "Baby Ruth"           "Boston Baked Beans"
[9] "Candy Corn"          "Caramel Apple Pops"
[11] "Charleston Chew"     "Chewey Lemonhead Fruit Mix"
[13] "Chiclets"            "Dots"
[15] "Dum Dums"            "Fruit Chews"
[17] "Fun Dip"             "Gobstopper"
[19] "Haribo Gold Bears"   "Haribo Happy Cola"
[21] "Haribo Sour Bears"   "Haribo Twin Snakes"
[23] "Hershey's Kisses"    "Hershey's Krackel"
[25] "Hershey's Milk Chocolate" "Hershey's Special Dark"
[27] "Jawbusters"          "Junior Mints"
[29] "Kit Kat"             "Laffy Taffy"
[31] "Lemonhead"           "Lifesavers big ring gummies"
[33] "Peanut butter M&M's" "M&M's"
[35] "Mike & Ike"           "Milk Duds"
[37] "Milky Way"           "Milky Way Midnight"
[39] "Milky Way Simply Caramel" "Mounds"
[41] "Mr Good Bar"         "Nerds"
[43] "Nestle Butterfinger" "Nestle Crunch"
[45] "Nik L Nip"           "Now & Later"
[47] "Payday"              "Peanut M&Ms"
[49] "Pixie Sticks"        "Pop Rocks"
[51] "Red vines"           "Reese's Miniatures"
[53] "Reese's Peanut Butter cup" "Reese's pieces"
```

```

[55] "Reese's stuffed with pieces" "Ring pop"
[57] "Rolo" "Root Beer Barrels"
[59] "Runts" "Sixlets"
[61] "Skittles original" "Skittles wildberry"
[63] "Nestle Smarties" "Smarties candy"
[65] "Snickers" "Snickers Crisper"
[67] "Sour Patch Kids" "Sour Patch Tricksters"
[69] "Starburst" "Strawberry bon bons"
[71] "Sugar Babies" "Sugar Daddy"
[73] "Super Bubble" "Swedish Fish"
[75] "Tootsie Pop" "Tootsie Roll Juniors"
[77] "Tootsie Roll Midgies" "Tootsie Roll Snack Bars"
[79] "Trolli Sour Bites" "Twix"
[81] "Twizzlers" "Warheads"
[83] "Welch's Fruit Snacks" "Werther's Original Caramel"
[85] "Whoppers"

```

```
candy["Milky Way", ]
```

	chocolate	fruity	caramel	peanut	almond	nougat	crisp	rice	wafer	hard
Milky Way	1	0	1	0	1				0	0
	bar	pluribus	sugar	percent	price	percent	win	percent		
Milky Way	1	0	0.604	0.651	73.09956					

```
candy["Sour Patch Kids",]$winpercent
```

```
[1] 59.864
```

```
candy["Dum Dums",]$winpercent
```

```
[1] 39.46056
```

```
skimr::skim(candy)
```

Table 1: Data summary

Name	candy
Number of rows	85
Number of columns	12
Column type frequency: numeric	12
Group variables	None

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalmondy	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedricewafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	


Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

winpercent

Q8. Plot a histogram of winpercent values


```
hist(candy$winpercent)
```

Histogram of candy\$winpercent



```
library(ggplot2)

ggplot(candy) +
  aes(winpercent) +
  geom_histogram(bins=10, col="red", fill="yellow")
```



Q11. On average is chocolate candy higher or lower ranked than fruit candy?

```
chocolate.inds <- as.logical(candy$chocolate)
chocolate.win <- candy[chocolate.inds,]$winpercent
mean(chocolate.win)
```

```
[1] 60.92153
```

And for fruit candy...

```
fruit.inds <- as.logical(candy$fruity)
fruit.win <- candy[fruit.inds,]$winpercent
mean(fruit.win)
```

```
[1] 44.11974
```

```
t.test(chocolate.win, fruit.win)
```

Welch Two Sample t-test

```
data: chocolate.win and fruit.win
t = 6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 11.44563 22.15795
sample estimates:
mean of x mean of y
 60.92153  44.11974
```

3. Overall Candy Rankings

The base R `sort()` and `order()` functions are very useful!

```
x <- c(5,1,2,6)

sort(x)
```

```
[1] 1 2 5 6
```

```
x[ order(x) ]
```

```
[1] 1 2 5 6
```

```
y <- c("barry", "alice", "chandra")
y
```

```
[1] "barry"  "alice"  "chandra"
```

```
sort(y)
```

```
[1] "alice"  "barry"  "chandra"
```

```
order(y)
```

```
[1] 2 1 3
```

First I want to order/arrange the whole dataset by winpercent values

```
inds <- order(candy$winpercent)

head(candy[inds,], n=5)
```

	chocolate	fruity	caramel	peanut	almond	nougat		
Nik L Nip	0	1	0		0	0		
Boston Baked Beans	0	0	0		1	0		
Chiclets	0	1	0		0	0		
Super Bubble	0	1	0		0	0		
Jawbusters	0	1	0		0	0		


	crisped	rice	wafer	hard	bar	pluribus	sugar	percent	price	percent
Nik L Nip		0	0	0		1		0.197		0.976
Boston Baked Beans		0	0	0		1		0.313		0.511
Chiclets		0	0	0		1		0.046		0.325
Super Bubble		0	0	0		0		0.162		0.116
Jawbusters		0	1	0		1		0.093		0.511

	winpercent
Nik L Nip	22.44534
Boston Baked Beans	23.41782
Chiclets	24.52499
Super Bubble	27.30386
Jawbusters	28.12744

Barplot

The default barplot, made with `geom_col()` has the bars in the order they are in the dataset...

```
ggplot(candy) +
  aes(winpercent, rownames(candy)) +
  geom_col()
```

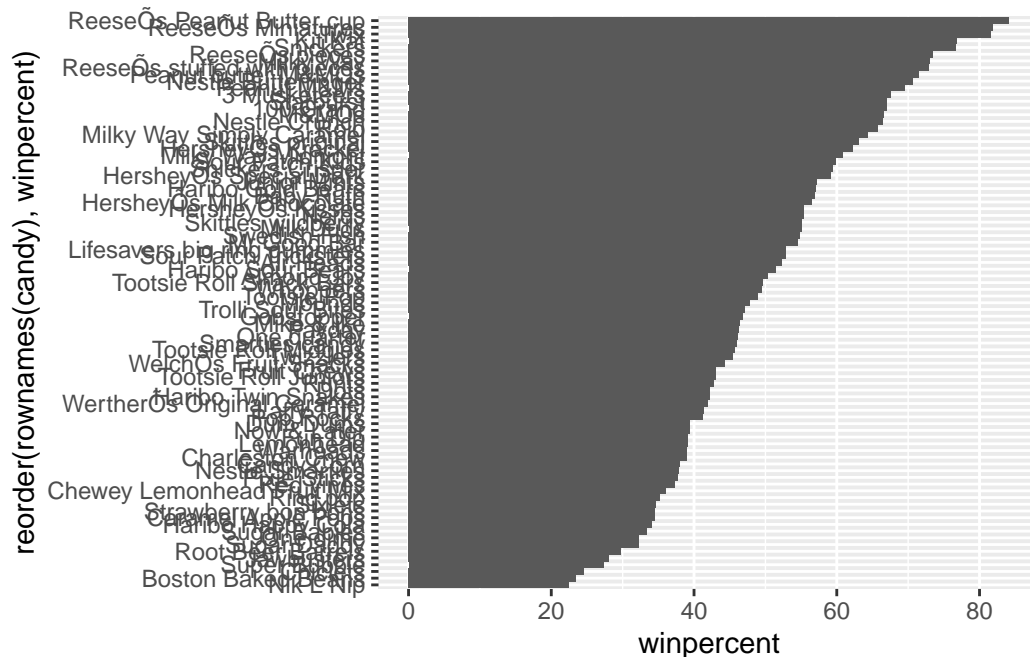



```
p <- ggplot(candy) +
  aes(winpercent, reorder(rownames(candy), winpercent)) +
  geom_col()
```

```
ggsave("mybarplot.png", p)
```

Saving 5.5 x 3.5 in image

p



Let's setup a color vector (that signifies candy type) that we can then use for some future plots. We start by making a vector of all black values (one for each candy). Then we overwrite chocolate (for chocolate candy), brown (for candy bars) and red (for fruity candy) values.


```
my_cols <- rep("black", nrow(candy))
#my_cols
my_cols[ as.logical(candy$chocolate) ] <- "chocolate"
my_cols[ as.logical(candy$bar) ] <- "brown"
my_cols[ as.logical(candy$fruity) ] <- "red"
my_cols
```

```
[1] "brown"      "brown"      "black"      "black"      "red"        "brown"
[7] "brown"      "black"      "black"      "red"        "brown"      "red"
[13] "red"        "red"        "red"        "red"        "red"        "red"
[19] "red"        "black"      "red"        "red"        "chocolate"  "brown"
[25] "brown"      "brown"      "red"        "chocolate"  "brown"      "red"
[31] "red"        "red"        "chocolate"  "chocolate"  "red"        "chocolate"
[37] "brown"      "brown"      "brown"      "brown"      "brown"      "red"
[43] "brown"      "brown"      "red"        "red"        "brown"      "chocolate"
[49] "black"      "red"        "red"        "chocolate"  "chocolate"  "chocolate"
[55] "chocolate"  "red"        "chocolate"  "black"      "red"        "chocolate"
[61] "red"        "red"        "chocolate"  "red"        "brown"      "brown"
```

```
[67] "red"      "red"      "red"      "red"      "black"    "black"
[73] "red"      "red"      "red"      "chocolate" "chocolate" "brown"
[79] "red"      "brown"    "red"      "red"      "red"      "black"
[85] "chocolate"
```

Now I can use this vector to color up my barplot

```
ggplot(candy) +
  aes(winpercent, reorder(rownames(candy), winpercent)) +
  geom_col(fill=my_cols)
```




4. Taking a look at pricepercent

What about value for money? What is the the best candy for the least money?

One way to get at this would be to make a plot of winpercent vs the pricepercent variable.

```
ggplot(candy) +
  aes(winpercent, pricepercent, label=rownames(candy)) +
  geom_point(col=my_cols) +
  geom_text(col=my_cols)
```




This plot sucks! I can not read the labels... We can use ggrepel package to help with this

```
library(ggrepel)

ggplot(candy) +
  aes(winpercent, pricepercent, label=rownames(candy)) +
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols, max.overlaps = 7)
```

Warning: ggrepel: 68 unlabeled data points (too many overlaps). Consider increasing max.overlaps

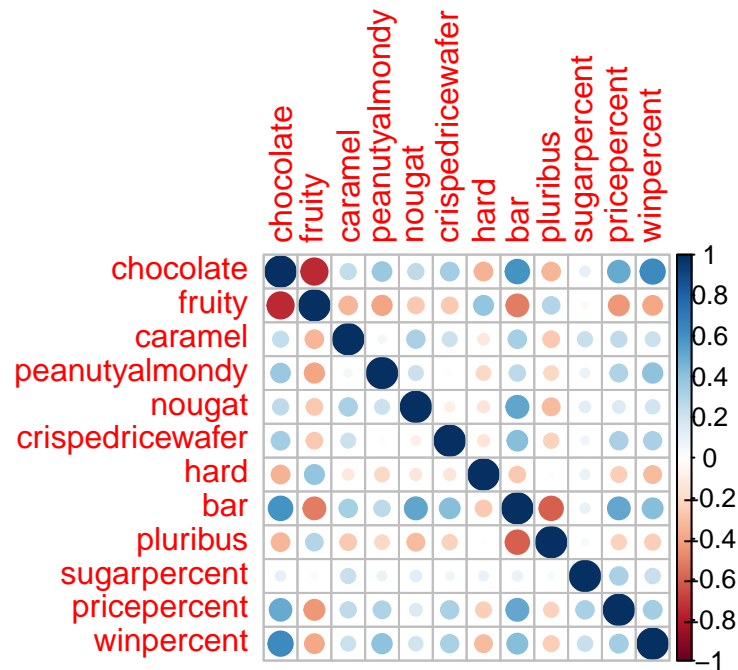


5. Exploring the correlation structure

```
library(corrplot)
```

corrplot 0.92 loaded

```
cij <- cor(candy)
corrplot(cij)
```



PCA: Principal Component Analysis

The main function that always there for us is `prcomp()`. It has an important argument that is set to `scale=FALSE`.

```
pca <- prcomp(candy, scale=TRUE)
summary(pca)
```


Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369

	PC8	PC9	PC10	PC11	PC12
Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317
Cumulative Proportion	0.89998	0.93832	0.97071	0.98683	1.00000

My PCA plot (a.k.a.) PC1 vs PC2 score plot.

```
plot(pca$x[,1], pca$x[,2], col=my_cols, pch=16)
```




I will make a “nicer” plot with ggplot. ggplot only works with data.frames as input so I need to make one for it first...

```
# Make a new data-frame with our PCA results and candy data
my_data <- cbind(candy, pca$x[,1:3])
```

```
ggplot(my_data) +
  aes(PC1, PC2, label=rownames(my_data)) +
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols, max.overlaps = 7)
```

Warning: ggrepel: 63 unlabeled data points (too many overlaps). Consider increasing max.overlaps



```
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530
Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369

	PC8	PC9	PC10	PC11	PC12
Standard deviation	0.74530	0.67824	0.62349	0.43974	0.39760
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317
Cumulative Proportion	0.89998	0.93832	0.97071	0.98683	1.00000